

Measuring E-Mail Header Injections on the World Wide Web*

Sai Prashanth Chandramouli[†], Pierre-Marie Bajan[‡], Christopher Kruegel[§]

Giovanni Vigna[§], Ziming Zhao[†], Adam Doupe[†], and Gail-Joon Ahn^{†,‡}

[†]Arizona State University, [‡]IRT SystemX, [§]University of California, Santa Barbara, [‡]Samsung Research

[†]{saipc, zzhao30, doupe, ahn}@asu.edu, [‡]pierre-marie.bajan@irt-systemx.fr, [§]{chris, vigna}@cs.ucsb.edu

ABSTRACT

E-mail header injection vulnerability is a class of vulnerability that can occur in web applications that use user input to construct e-mail messages. E-mail header injection vulnerabilities exist in the built-in e-mail functionality of the popular languages PHP, Java, Python, and Ruby. With the proper injection string, this vulnerability can be exploited to allow an attacker to inject additional headers, modify existing headers, and alter the content of the e-mail.

While E-mail header injection vulnerabilities are known to the community, and some commercial vulnerability scanners claim to discover E-mail header injection vulnerabilities, they have never been studied by the academic community. This paper presents a scalable mechanism to automatically detect E-mail header injection vulnerabilities and uses this mechanism to quantify the prevalence of E-mail header injection vulnerabilities on the web. From crawling 23,553,796 URLs, we found 994 vulnerable URLs across 414 domains. 135 of these domains are in the Alexa top 1 million, and five of them are in the top 20,000. 137 of the vulnerable domains are using anti-spoofing mechanisms such as DKIM, SPF, or DMARC, and E-mail header injection renders this protection useless. This work shows that E-mail header injection vulnerabilities are widespread and deserve future research attention.

CCS CONCEPTS

• Security and privacy → Software and application security; Web application security;

KEYWORDS

E-mail header injection, Software Security

1 INTRODUCTION

The World Wide Web has single-handedly brought about a change in the way we use computers. The ubiquitous nature of the web makes it possible for anyone to access information and services anywhere and on multiple devices such as phones, laptops, TVs, and cars. This access has ushered in an era of web applications which depend on user input. While the rapid pace of development has improved the speed of information dissemination, it comes at a cost. As users move more and more of their personal and financial

information to web applications, attackers are responding by using web application vulnerabilities to steal lucrative data.

Many common and well-known web application vulnerabilities, such as SQL Injection and Cross-Site Scripting [41], are command injection vulnerabilities [55], where malicious user input is used to alter the structure of a command (SQL query and JavaScript code respectively). Developers of web applications must use the correct sanitization routine in all paths from user input to a command.

E-mail header injection vulnerabilities are a lesser-known command injection vulnerability. We verified that this vulnerability exists in the implementation of the built-in e-mail functionality in the popular languages PHP, Java, Python, and Ruby. The format of e-mail messages is defined by the Simple Mail Transfer Protocol (SMTP) [48]. Each e-mail message is represented by a series of headers separated by newlines, followed by the body content (separated from the headers by two newlines). Some of these headers are mandatory (From, To, Date), but the headers could also include other information such as the Subject, BCC, etc.

With the proper exploit injection string, E-mail header injection vulnerabilities can be exploited by an attacker to inject additional headers, modify existing headers, or alter the contents of the e-mail—while still appearing to be from a legitimate source. E-mail header injection exploits allow an attacker to perform e-mail spoofing, resulting in phishing attacks *that are sent from the actual e-mail server*, and thus bypass e-mail spoofing technologies, such as DKIM [12], SPF [51], and DMARC [37].

While some command injection vulnerabilities have received extensive attention from the research community, E-mail header injection vulnerabilities have received little focus. In fact, the Acunetix vulnerability scanner contains an AcuMonitor component which claims to detect E-mail header injection vulnerabilities while scanning [10]. Unfortunately, as a commercial product, little is known about how AcuMonitor detects E-mail header injection vulnerabilities.

To shed light on this little-studied vulnerability class, we describe E-mail header injection vulnerabilities and measure E-mail header injection vulnerabilities. To perform this measurement, we crawled the web, extracted forms with e-mail fields, and injected them with different payloads to infer the existence of an E-mail header injection vulnerability. We then automatically audited received e-mails to see if any of the injected data was present. This allowed us to classify whether a particular URL was vulnerable to the attack. Our automated system works in a black-box manner, without looking at the web application's source code, and only analyzes the payloads in the e-mails.

In summary, we make the following contributions:

- We develop a black-box approach to detect E-mail header injection vulnerabilities in a web application.
- We develop an open-source system to crawl the web and automatically detect E-mail header injection vulnerabilities.

*This work was partially supported by a grant from the Center for Cybersecurity and Digital Forensics at Arizona State University

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

SAC 2018, April 9–13, 2018, Pau, France

© 2018 Association for Computing Machinery.

ACM ISBN 978-1-4503-5191-1/18/04...\$15.00

<https://doi.org/10.1145/3167132.3167308>

```

1 $from = $_REQUEST['email'];
2 $subject = 'Hello XYZ';
3 $message = 'We need you to reset your password';
4 $to = 'xyz@example.com';
5 $returnValue = mail($to, $subject, $message, "From: $from");

```

Listing 1: PHP program with e-mail header injection vulnerability.

CVE No.	Affected Software	Year
2002-1575	cgiemail	2004
2002-1771	FormMail 1.9	2005
2002-1917	Geeklog 1.35	2005
2005-0493	Biz Mail Form <=2.2	2005
2005-2854	thesitewizard.com	2005
2005-3883	PHP mb_send_mail	2005
2006-0631	Perl mailback.pl	2006
2006-0712	Squishdot 1.5.0	2006
2006-1225	Drupal 4.5.0-4.5.8 and 4.6.0-4.6.8	2006
2006-1305	Microsoft Outlook 2000, 2002-03	2006
2006-2159	Russcom Network	2006
2006-2943	CGI-RESCUE WebFORM 4.1	2006
2006-2944	CGI-RESCUE FORM2MAIL 1.21	2006
2006-3171	CS-Forum <=0.82	2006
2006-3473	Drupal Module <=1.8.2.2	2006
2006-4344	CGI-Rescue Mail	2006
2006-7020	phpwms 1.2.5-DEV	2007
2006-7087	Dotdeb PHP	2007
2007-1718	PHP 4.0-4.4.6 and 5.0-5.2.1	2007
2007-1898	Jetbox CMS 2.1	2007
2007-1900	FILTER_VALIDATE_EMAIL PHP	2007
2007-2731	Jetbox CMS 2.1	2007
2008-2105	Bugzilla	2008
2009-1469	IceWarp	2009
2008-7281	OTRS - Open Ticket Request System	2011
2014-2957	Exim	2014
2015-8476	PHPMailer	2015
2016-4803	dotCMS	2016

Table 1: History of software found in Common Vulnerabilities and Exposures database affected by e-mail header injection vulnerability

- We use our system to crawl 23,553,796 URLs, and we find 994 URLs vulnerable to E-mail header injection across 414 domains.
- We perform detailed analyses on the domains found to be vulnerable: finding their Alexa rankings, the technologies used in such vulnerable domains, presence of e-mail spoofing protection, presence of such domains on existing spam-lists, and the ability to send malicious attachments.

2 BACKGROUND

E-mail header injection belongs to the class of command injection vulnerabilities. However, unlike its more popular siblings, SQL injection [18, 27, 50], Cross-Site Scripting [32, 35], or HTTP Header Injection [33], relatively little research is available on E-mail header injection vulnerabilities.

As with other command injection vulnerabilities, E-mail header injection is caused by improper or nonexistent sanitization of user input. If the program constructs e-mails from user input and fails to check for the presence of e-mail headers, a malicious user can control the headers of this particular e-mail. E-mail header injection vulnerabilities can be leveraged to enable malicious attacks, including, but not limited to, spoofing or phishing.

2.1 History of E-mail header injection

We found the first E-mail header injection description in a late 2004 article on phpsecure.info [58] accredited to user tobozo describing how an E-mail header injection vulnerability existed in the implementation of the mail() function in PHP and how it can

```

1 Received: from mail.ourdomain.com ([62.121.130.29])
2   by xyz.com (Postfix) with ESMTP id 5A08E52C0154
3   for <abc@example.com>; Sun, 20 Mar 2016 13:56:58 -0700 (MST)
4 From: abc@example.com
5 To: xyz@example.com
6 Subject: Hello XYZ
7 CC: spc@example.com
8 Date: Sun, 20 Mar 2016 13:56:58 -0700(MST)
9
10 We need you to reset your password

```

Listing 2: SMTP headers generated by a PHP mail script.

be exploited. More recently, a blog post by Damon Kohler [36] and an accompanying wiki article [52] describe the attack vector and outline few defense measures for E-mail header injection vulnerabilities.

An example of the vulnerable code written in PHP is shown in Listing 1. This code takes in user input from \$_REQUEST['email'], and stores it in the variable \$from, which is later passed to the mail() function to construct and send the e-mail.

When this code is given the malicious input abc@example.com\nCC:spc@example.com as the value of the \$_REQUEST['email'], it generates the equivalent SMTP Headers shown in Listing 2. It can be seen that the CC (carbon copy) header that we injected appears as part of the resulting SMTP message. This will make the e-mail get sent to the e-mail address specified as part of the CC as well.

We gathered reported Common Vulnerabilities and Exposures (CVE) [4] to get an idea of the distribution of reported E-mail header injection vulnerabilities over time. From the 28 reports we found (Table 1), it can be seen that even though many vulnerabilities were found in earlier years (2005-07), there have been recently discovered E-mail header injection vulnerabilities which suggests that it is still a very real and relevant threat to modern web security.

2.2 Languages Affected

PHP was the first language found vulnerable to E-mail header injection in its implementation of the mail() function at the time of release of PHP 4.0. According to w3techs [59], PHP is used by 81.9% of all websites.

After 13 further iterations of PHP since the 4.0 release (the current version is 7.1), the mail() function is yet to be fixed after 15 years. However, the PHP documentation [43] specifies that the mail() function does not protect against E-mail header injection. A working code sample with the vulnerability is shown in Listing 1.

A bug was filed about an E-mail header injection vulnerability in Python's implementation of the email.header library and the header parsing functions in early 2009, which was followed by a partial patch in 2011.

Unfortunately, the bug fix was only for the email.header package, and not for other frequently used packages such as email.parser, where both the classic Parser() and the newer FeedParser() contain E-mail header injection vulnerabilities even in the latest versions: 2.7.11 and 3.5. The bug fix was also not backported to older versions of Python. There is no mention of the vulnerability in the Python documentation for either library. Contrary to PHP's behavior of overwriting existing headers, Python only recognizes the first occurrence of a header, and ignores duplicate headers.

Java has a bug report about E-mail header injection filed against its JavaMail API. A detailed write-up by Alexandre Herzog [28] contains a proof-of-concept program that exploits the API to inject headers.

```

1 Received: from mail.ourdomain.com ([62.121.130.29])
2   by xyz.com (Postfix) with ESMTP id 5A08E52C0154
3   for <abc@example.com>; Sun, 20 Mar 2016 13:56:58 -0700 (MST)
4 From: abc@example.com
5 CC: 1@example.com, 2@example.com, 3@example.com
6 Subject: My Subject
7 Content-Type: multipart/mixed; boundary=foobar;
8 --foobar
9 Content-Type: text/html
10
11 This is the attacker's body
12 --foobar
13 To: xyz@example.com
14 Subject: Hello XYZ
15 Date: Sun, 20 Mar 2016 13:56:58 -0700(MST)
16
17 We need you to reset your password

```

Listing 3: Exploiting the E-mail header injection vulnerability in Listing 1 to control the recipients, subject, and body of the SMTP message.

From our testing, Ruby’s built-in Net: :SMTP library also has the vulnerability (not documented on the library’s homepage).

2.3 Exploitation

Successful exploitation of an E-mail header injection vulnerability depends on where injection occurs in the SMTP message. The attacker cannot alter parts of the SMTP message that precede the injection location, but the attacker has complete control over everything that follows. However, similar to other command injection vulnerabilities, the remaining parts of the SMTP message will always be appended to the attacker’s injection, so the attacker must contend with this. By exploiting an E-mail header injection vulnerability, an attacker can control who receives the message (and can include multiple CC and BCC recipients), the body, and possibly the subject (depending on if the subject header occurs before/after the injection point and the language used).

The exploitation techniques that we describe below were discussed in our prior work [20]. The main vector for exploiting E-mail header injection vulnerabilities follows the template of command injection vulnerability exploitation: first inject the attacker’s desired commands, then comment out the rest of the message. In E-mail header injection vulnerabilities, the attacker first includes all SMTP headers she desires. These will typically be the Subject header to control the subject of the e-mail¹, CC or BCC headers to control the recipients of the e-mail.

To handle the extra content after the injection point, one technique is to use a Content-type header to specify that the SMTP message is a multi-part email and that the sections are separated by an attacker-specified boundary. The boundary delineates different parts of the message so that the attacker’s body is the only valid part of the message, and the attacker can choose a random value for the boundary that is not present in the rest of the SMTP message.

Using this technique, the attacker can control the e-mail. Injecting the following attack payload: “abc@example.com\nCC:1@example.com, 2@example.com\nSubject: Pwned\nContent-Type: multipart/mixed;boundary=foobar;\n--foobar\nContent-Type: text/html\n\nThis is the attacker's body\n--foobar” into the email parameter of the PHP program in Listing 1 results in the SMTP message shown in Listing 3. By expanding on this technique, the attacker can include links in the e-mail, or even attachments, by adding additional multipart messages with different content types.

¹The SMTP protocol specifies that there should only be one Subject header, so the attacker may not be able to alter the subject if the header is already defined. This behavior would be MUA-dependent.

A shorter technique, if the injection point is limited in input size, is to use an HTML comment to ignore the developer-controlled part of the SMTP message, using a payload like: “abc@example.com\nCC:1@example.com, 2@example.com\nSubject: Random\nContent-Type: text/html\n\nThis is the attacker's body<!--”. However, this technique will only work if the developer-controlled part of the SMTP message does not contain a closing HTML comment tag -->.

2.4 Impact of E-mail header injection

The impact of an E-mail header injection vulnerability can be far-reaching. According to w3tech, PHP, Java, Python, and Ruby (combined) account for over 85% of the server-side programming languages of websites measured, and the default implementation of the e-mail functionality in these languages is vulnerable to E-mail header injection. An E-mail header injection vulnerability can be exploited to do potentially any of the following:

Phishing and Spoofing Attacks Phishing [31] (a variation of spoofing [24]) refers to an attack where the recipient of an e-mail is made to believe that the e-mail is legitimate when it was really created by a malicious party. The e-mail usually redirects the victim to a malicious website, which then steals their credentials or infects their computer with malware (via a drive-by-download).

E-mail header injection gives attackers the ability to inject arbitrary headers into an e-mail sent by a website *and control the output of the e-mail*. This adds credibility to the generated e-mail, as it is sent from the website’s mail server and users (and anti-spam defenses) are more likely to trust an e-mail that is received from the proper mail server. Therefore, attackers could leverage E-mail header injection vulnerabilities to perform enhanced phishing attacks.

Spam Networks Spam networks can use E-mail header injection vulnerabilities to send a large amount of e-mail from servers that are trusted. By adding additional cc or bcc headers to the generated e-mail, attackers can easily choose the recipient of the spam email.

Due to the e-mail being from trusted domains, recipient e-mail clients and anti-spam systems might not flag them as spam. If they do flag them as spam, then that can lead to the website being blacklisted as a spam generator (which would cause a Denial of Service on the vulnerable web application).

Information Extraction E-mails can contain sensitive data that is meant to be accessed only by the user. Due to an E-mail header injection vulnerability, an attacker can add a bcc header, and the e-mail server will send a copy of the e-mail to the attacker, thereby exposing important information. User privacy can thus be compromised, and loss of private information can lead to other attacks.

Denial of Service Denial of service attacks (DoS), can also be caused by exploiting an E-mail header injection vulnerability to send excessive e-mails resulting in overloading the mail server and cause crashes or instability.

3 SYSTEM DESIGN

To quantify the existence of E-mail header injection vulnerabilities on the web at large, we developed a system to automatically detect E-mail header injection vulnerabilities in a black-box manner.

3.1 Approach

We took a black-box approach to measure the prevalence of E-mail header injection vulnerabilities on the web. Black-box testing [15]

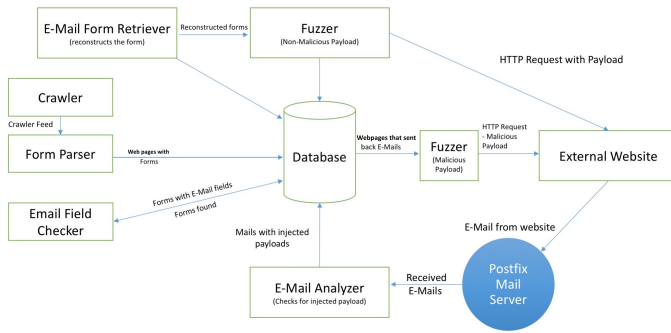


Figure 1: Overall system architecture.

is a way to examine the functionality of an application without analyzing its source code. Black-box testing allows our system to detect E-mail header injection vulnerabilities in *any* server-side language (not simply those we identified in Section 2.2). The overall architecture of our system is presented in Figure 1.

3.2 System Components

The Data Gathering module and Payload Injection modules are composed of smaller components. This section describes the functionality of those components.

3.2.1 Data Gathering Module. We used an open-source Apache Nutch based Crawler [1]. The Crawler provides the system with a continuous feed of URLs and the HTML contained in those pages. The Form Parser is responsible for parsing the HTML and retrieving data about the HTML forms on the page, including the following: (1) Form attributes, such as method and action (URL) for the HTTP request, (2) Data about form inputs, such as their attributes, names, and default values. The default values are essential for fields like `<input type="hidden">` as these fields are usually used to check for the submission of forms by bots, and (3) Presence of the `<base>` element in the HTML, as this affects the final URL to which the form is to be submitted (if the action attribute is a relative URL).

The E-mail Field Checker is the final stage in the Data Gathering module. It receives the HTML form data and checks for the presence of e-mail fields in those forms. If any e-mail fields are found, it stores references to these forms. The intuition here is that we do not want to try to fuzz all HTML forms on the web to look for E-mail header injection vulnerabilities, rather just those HTML forms that are likely to invoke server-side email functionality.

The E-mail Field Checker searches for the words e-mail, mail or email within the form, instead of an explicit HTML5 e-mail field (e.g., `<input type="email">`). This is by design, taking into account a common design pattern used by web developers, where they may have a text field with an id or name attribute set to email, instead of an actual e-mail type attribute, for purposes of backward compatibility with older browsers. The output of this stage is stored in the database and acts as the input to the Payload Injection module.

3.2.2 Payload Injection Module. E-mail Form Retriever is the first stage in the Payload Injection module. It does the following: (1) Retrieve forms and remove duplicates, (2) reconstruct each form's input fields and values from stored form data, and (3) construct the target URL to create an HTTP request for fuzzing.

The Fuzzer interacts directly with the external web applications. It injects payloads in two stages to reduce the total number of HTTP requests the system generates to detect an E-mail header injection vulnerability. Making HTTP requests is an expensive process [17], and can cause bottlenecks in a Crawler-Fuzzer system [53]. The two different payloads used for fuzzing are:

Non-Malicious Payload. The non-malicious payload is simply an e-mail address. The goal is to see if the web application will send an e-mail message based on our input. The specific format of the e-mail is `reguser#@example.com`, where `\#` is replaced by an internal id that uniquely maps the payload to the form, and `example.com` is replaced by our domain.

Malicious Payload. After receiving an e-mail from a specific form, we use the malicious payload to try to exploit an E-mail header injection vulnerability. We inject the form fields with the bcc (blind carbon copy) header. If the vulnerability is present, this will cause the server to send a copy of the e-mail to the e-mail address we added in the bcc field.

We consider a special case: the addition of an x-check header field to the payloads. This is due to Python's exhibited behavior when attaching headers. Instead of overwriting a header if it is already present, Python will ignore duplicate headers. So, if the bcc field is already present as part of the headers, our injected bcc header would be ignored. To overcome this, we inject a new header that is not likely to be generated by the web application.

We created four different malicious payloads. Each of these payloads is crafted for a particular use case. The four payloads are: (1) `user#@example.com\nbcc:maluser#@example.com`, (2) `user#@example.com\nbcc:maluser#@example.com\nx-check:in`, (3) `user#@example.com\r\nbcc:maluser#@example.com`, and (4) `user#@example.com\r\nbcc:maluser#@example.com\r\nx-check:in`.

Payload 1 is the most minimal payload: it injects a newline character followed by the bcc field. Payload 2 contains the additional x-check header to inject Python-based web applications. Payloads 3 and 4 are added for purposes of cross-platform fuzzing: `\r\n` is the “Carriage Return - Line Feed (CRLF)” used on Windows systems [25]. The `#` are replaced by an id, for mapping to the forms.

Along with the payload, the Fuzzer also injects data into the other fields of the form. This data must pass validation constraints on the individual input fields (e.g., name field might not allow numbers). As crawling and fuzzing input fields on the web is an open problem [47], we chose to go with a best-effort approach. To maximize the amount of vulnerabilities the system discovers, the data injected into the input fields should adhere to the constraints. The Fuzzer uses a “Data Dictionary” which has predefined “keys” and “values” for standard input fields such as name, date, username, password, text, and submit. The values for these are generated for each form, based on generally followed guidelines for such fields. For example, password fields should consist of at least one uppercase letter, one lowercase letter, and special characters.

When the fuzzed data is ready, the Fuzzer constructs the appropriate HTTP request (GET or POST) and sends the HTTP request to the URL that was generated by the E-mail Form Retriever.

Injection Verification stage checks for the presence of injected data in the received e-mails. This stage operates on the e-mails received and stored by our Postfix server, and, depending on the user account that received the e-mail, it performs different functions.

Analyzing regular e-mail. ‘Regular e-mail’ refers to the e-mails received by account reguser#@example.com that were sent due to

Type of Data	Quantity
URLs Crawled	23,553,796
Total Forms found	7,354,425
Forms with E-Mail Fields	1,228,774

Table 2: The data collected for our project.

Type of fuzzing	Forms fuzzed	E-Mails received
Regular payload	1,012,530	74,192
Malicious payload	64,510	994

Table 3: The data we fuzzed and the e-mails we received.

injecting the regular, non-malicious payload. The objective of the analysis on this e-mail is identify if the input fields that we injected with data appear on the resulting e-mail, and if so, which fields appear where.

To find this, we parse each received e-mail, and check whether any of the fields we injected with data appear in either the headers or body of the e-mail. These could be fields such as name, user-name, age, etc. If they do, we add them to the list of fields that can *potentially* result in an E-mail header injection vulnerability for the given e-mail.

We then pass on this information back to the Fuzzer pipeline, along with the vulnerable form, where these fields are *also* fuzzed along with the e-mail fields to check for the presence of E-mail header injection.

Analyzing e-mail with payloads The “e-mails with payloads” refers to e-mails received by either the `nuser#@example.com` or `maluser#@example.com` accounts. These e-mails could only be received as a result of injecting the malicious payloads.

Detecting injected bcc headers The structure of our payload ensures that any e-mails received directly by the `maluser` account indicate the presence of the injected bcc field.

Detecting injected x-check headers E-mails not received by the `maluser` account but by the `nuser` account constitute a special category of e-mails. These e-mails could have been generated due to two reasons: (1) The web application performed sanitization and stripped out the bcc part of the payload, thereby sending e-mails only to the `nuser` account. These e-mails then act as proof that the vulnerability was not found on the given URL. (2) The bcc header can be ignored for other reasons (e.g., Python’s default behavior when it encounters duplicate headers). In this case, we check if the e-mail contains the custom header `x-check`. If it does, then this is a successful exploit of the vulnerability.

4 EVALUATION

We ran our system on the web at large, attempting to discover E-mail header injection vulnerabilities in web applications. We used the Alexa top 10,000 websites as well as a feed of 10,000 random blog pings per day from `weblogs.com` as a seed to `Crawler`. All domains were crawled to a maximum depth of four, and the crawler respected the `robots.txt` directive. This seed enabled `Crawler` to get an overview of not only the popular websites (from the Alexa top 10k) but also the long-tail of blog posts and websites that they linked to.

From our extensive crawl of the web, we were able to gather the data shown in Table 2. We ran the system for 76 days, during which our system crawled 23,553,796 unique URLs, and found a total of 7,354,425 forms from 1,085,365 unique domains. Out of these forms, our system found 1,228,774 forms that contained an e-mail field, from 198,306 unique domains. Table 3 shows the quantity of e-mails we received for the benign and malicious payloads.

Type of Injection	# e-mails
E-Mail Header Injections with bcc header	583
E-Mail Header Injections with x-check header	493
To header injections alone	229
Injections with both bcc and x-check headers	310
Both To header injections and x-check headers	15
x-check headers found in nuser e-mails	239
Unique x-check headers found in nuser e-mails	182
Total successful injections (1 + 3 + 7)	994

Table 4: Classification of the e-mails that we received into broad categories of the vulnerability.

E-mail received from forms. The e-mails that we received can be categorized into two categories. (1) E-mails due to regular payload: This represents the total number of web applications that sent e-mails to us. This indicates that we were able to successfully submit the forms on these sites to trigger the web application to send an e-mail. (2) E-mails due to malicious payload: Once we receive an e-mail from a web application due to the regular payload, we fuzz those forms with the malicious payloads. This field represents the total number of unique URLs that contain an E-mail header injection vulnerability.

4.1 Analysis of the Received E-mail Data

During our analysis of the received e-mails, we found that the e-mails that we received belonged to three categories:

(1) E-mails with the bcc header successfully injected. This form of injection was our initial objective, and we found 583 such e-mails in our received e-mails. This validates that the web applications that sent these e-mails are vulnerable to E-mail header injection.

(2) E-mails with the To header successfully injected. We discovered an unintended vulnerability class during our analysis, which we call *To header injection*. These injections reflect the ability to inject any number of e-mail addresses into the `To` field of the SMTP message while being unable to inject any other header into the e-mails. We found 229 such e-mails in our received e-mails. We attribute this behavior to inconsistent sanitization by the application.

While not allowing us complete control over content of the e-mails sent, *To header injection* makes it possible to append any number of e-mail addresses, thereby enabling us to leak information or perform DoS (Denial of Service) attacks against the web application.

(3) E-mails with the x-check header successfully injected. The third category of e-mails received were e-mails with the x-check header injected. As discussed in Section 3.2.2, we can differentiate between unsuccessful attempts and successful attempts by injecting the additional header and checking whether headers other than the bcc header can be injected into the generated e-mail. 493 e-mails were received with the x-check header injected.

We list each category and the number of e-mails received by that category in Table 4. We explain the combination of these header injections (4-7) as follows:

E-mail Header Injections with both bcc and x-check headers: These represent the scenario where an attacker can inject multiple headers into the e-mails. We can see that 54% of the received bcc header injected e-mails are also susceptible to being injected with additional headers.

Both To header injections and x-check headers: This combination shows us that in addition to being able to inject into the `To` fields, we injected additional headers into the e-mail. It is not clear what

Pipeline Stage	Quantity	Differential
Crawled URLs	23,553,796	$\Delta d2/d1 * 100$
Forms found	7,354,425	31.22%
E-Mail Forms found	1,228,774	16.71%
Fuzzed with regular payload	1,012,530	82.40%
Received e-mails	74,192	7.33%
Fuzzed with malicious payload	64,510	86.95%
Successful attacks	994	1.54%

Table 5: Data gathered by our pipeline at each stage

causes this behavior; however, these can be exploited to achieve the same result as a regular E-mail header injection.

Total x-check headers and unique x-check headers found in nuser e-mails: We found a total of 239 e-mails in the nuser account. Out of these, 182 had unique form ids that were *not* already found in the maluser account. We attribute these e-mails to (probably) being sent by a web application that was built with Python or another language having a similar behavior with respect to ignoring duplicate headers while constructing an e-mail, thus appending the x-check header and *not* the bcc header.

Total successful injections: This represents the total number of successful injections. This includes the E-mail header injection with bcc header (1), To header injections (3), and Unique x-check headers found in nuser e-mails (7). A total of 994 vulnerabilities were found by our system.

4.2 Understanding the Data Pipeline

Table 5 shows the data gathered by our pipeline, with the differential changes at each stage of the pipeline. At each stage of the pipeline, the amount of data decreases, for instance, out of the 23,553,796 URLs we crawled, only 7,354,425 forms (31.22%) were found. Out of these, only 1,228,774 forms (16.71%) contained e-mail fields.

In our fuzzing attempts, the same behavior is observed. We fuzzed 1,012,530 forms with the regular payload, which resulted in a total of 74,192 e-mails (7.33%). After analysis of the received e-mails, we further fuzzed 64,510 forms, which resulted in 994 e-mails (1.54%) which contain the vulnerability across 604 IP addresses from 414 domains.

We attribute the difference in the number of forms found to the number of forms fuzzed (a difference of 9,682 forms) to the presence of bot-blocking mechanisms on a website (discussed in Section 5.2), though we do not know what percentage was caused by each bot-blocking mechanism discussed in Section 5.2.

Note that over 1% of the forms that were not fuzzed (100 out of 9,682) were also tested manually using PostMan to generate HTTP requests with payloads to verify that our system was working as intended.

4.3 Analysis of Vulnerable domains

We performed the following analyses on the domains that were found to contain E-mail header injection vulnerabilities: (1) checking Alexa rankings of the vulnerable domains, (2) investigating the back-end languages used by the vulnerable domains, and (3) analyzing the presence of anti-spoofing defenses, such as DKIM [12], SPF [51], and DMARC [37].

4.3.1 Alexa rankings of vulnerable domains. We searched through the Alexa rankings data[7] for the domains that were found to be vulnerable, and found 135 of these domains in the top 1 million ranked websites, with 1 website in the top 5,000, and 4 in




Figure 2: Distribution of vulnerable domains based on Alexa Rankings

Technologies used	% of Vulnerable Domains
PHP	79
ASP.net	14
J2EE	5
Perl	1
Ruby on Rails	1
Wordpress	17

Table 6: Distribution of vulnerable domains based on technologies used

the top 20,000. A detailed distribution of the rankings is shown in Figure 2.

4.3.2 Back-end technologies of vulnerable domains. We investigated the top 100 vulnerable domains on our list (based on the Alexa rankings) to find the back-end technologies used by these vulnerable domains to see if there was any recurring pattern. Using BuiltWith [8] and Wappalyzer [9], we found that a majority of the vulnerable domains (79%) used PHP as one of the back-end technologies, while 17% of domains used WordPress and another 14% used ASP.Net. Other technologies used include Java (5%), Ruby on Rails (1%), and Perl (1%). We also found a combination of other technologies like Magento, CakePHP, CodeIgniter, Joomla, mail.ru and Drupal being used in conjunction with one of the above languages. 4% of the websites also used Contact Form 7 [2], which is supposed to protect against such attacks.

4.3.3 Presence of e-mail spoofing protection on vulnerable domains. We investigated the presence of anti-spoofing mechanisms such as DKIM (DomainKeys Identified Mail), SPF (Sender Policy Framework), and DMARC (Domain-based Message Authentication, Reporting & Conformance) in the vulnerable domains. Protection mechanisms such as these make it harder for an attacker to spoof the e-mail messages from these domains manually by forging the From addresses of e-mails.

Out of the top 100 vulnerable domains on our list (based on the Alexa rankings), we found that SPF protection was present on 71 of these websites, DKIM on 59 of these websites, and DMARC on 7 of these websites.

This makes E-mail header injection attacks all the more effective as now we have a way to spoof messages from domains that have anti-spoofing mechanisms in place preventing manual spoofing.

4.4 Exploitation Evidence

We compared the 604 IPs that our system found to be vulnerable to E-mail header injection against 13 well-known IP blacklists, to see if these IPs were being exploited by attackers to send spam. The blacklists that we used were: zen.spamhaus.org, spam.abuse.ch, cbl.abuseat.org, virbl.dnsbl.bit.nl, dnsbl.inps.de, ix.dnsbl.manitu.net, dnsbl.sorbs.net, bl.spamcannibal.org, bl.spamcop.net, dnsbl-1.uceprotect.net, dnsbl-2.uceprotect.net, dnsbl-3.uceprotect.net, db.wpbl.info.

We found that 157 of these IPs were blacklisted on at least one of the above blacklists for sending out spam, and 46 of them were found on multiple blacklists. We do not have enough data to make an observation about whether these attackers are exploiting E-mail header injection to send out the spam, as an alternative hypothesis is that these IPs are on the blacklists because the server has different vulnerabilities that attackers exploit to cause the server to send spam (assuming that the server is normally benign).

4.5 Emails with Malicious Attachments

As additional analysis to find domains on the internet that send out malicious attachments, we checked the e-mails received by the ‘reguser’ account (which are injected with regular e-mail addresses and not malicious data) for the presence of attachments that may contain malicious software, which will indicate that the server itself is not benign.

To do this, we passed the 2,950 attachments we received to Virus-Total [6] – an online malware scanner that checks for the presence of malware by running 40+ virus scanners on the uploaded files. We found that out of the 2,950 attachments, 443 were malicious, out of which 265 were from unique domains. This data is just a by-product of our research and could very well lead to future research.

4.6 Responsible Disclosure of Discovered Vulnerabilities

After we discovered an E-mail header injection vulnerability on a particular website, we attempted to notify the developers of the vulnerable web application, along with a brief description of the vulnerability. We chose to e-mail the following mailboxes, following the rules specified in RFC 2142 [21]: security@domain.com, admin@domain.com, and webmaster@domain.com

Out of the 414 vulnerable domains found, only 113 websites had the mailboxes able to receive e-mails. For the remaining domains, we used the whois [5] data to find the contact details of the owner and then e-mailed them. We received 21 developer responses, confirming 15 discovered vulnerabilities. Four of the developers fixed the vulnerability on their website.

From our research, it is clear that E-mail header injection is quite widespread as a vulnerability, appearing on 1.54% of forms that we were able to perform automated attacks on. This value acts as a *lower bound* for prevalence of E-mail header injection vulnerability, and can quite easily be larger if the attacks were broader, crafted for the individual web application, and less automated.

5 DISCUSSION

In this section, we discuss the lessons learned, the limitations of our system, and how to mitigate E-mail header injection vulnerabilities.

5.1 Lessons Learned

From our results, it is evident that E-mail header injection vulnerabilities exist in the wild. Despite its relatively low occurrence rate compared to the more popular SQL Injection and XSS (Cross-Site Scripting), when we consider total number of domains on the World Wide Web—1,018,863,952 according to Internet Live Stats [30] as of early 2016—and calculate 0.038% percent (the occurrence rate of E-mail header injection vulnerability calculated from vulnerable domains as found by our system to total number of domains crawled) of that number, this yields 295,693 domains. Of course, extrapolation in this way is not an accurate measure of the prevalence of E-mail header injection vulnerabilities. However, even with as few as a thousand domains affected by this vulnerability, it can still have a disastrous impact on these domains, and also on the overall World Wide Web due to the traffic caused by the sheer number of generated e-mails.

We found two different forms of E-mail header injection: the first one is the traditional one, injecting any header into the e-mail that allows the attacker complete control over the contents of the e-mail. The second attack has not yet been documented and provides the ability to inject multiple e-mail addresses into the To field. We call this a To header injection. In this vulnerability, an attacker can add addresses to the To field of the email with newlines separating the e-mail addresses. We could not determine if this vulnerability is due to unique flaws in each web application or if this vulnerability is due to an implementation issue with a particular language or framework. However, from our preliminary analysis, it is evident that the vulnerable web applications do not share much in common.

To header injection allows an attacker to extract information that should be private, and in some of these cases, able to inject enough data to spoof other headers of the e-mail message. From Table 4, information leakage using To header injection was possible on 229 forms, while spoofing using To header injection was possible on 15 forms.

5.2 Limitations

Because our system is fully automated, it is also susceptible to being stopped by mechanisms in web applications that prevent automated crawls or form submissions. A common reason for our fuzzing attempts to fail is the bot-blocking mechanisms built into the web applications. CAPTCHAs (Completely Automated Public Turing test to tell Computers and Humans Apart) [11] pose a very difficult problem for our system to exploit E-mail header injection, even if it is present. Other measures such as hidden form fields and CSRF (Cross-site Request Forgery [38]) tokens are also often used to detect bots [11, 46].

We made sure that we do not fuzz hidden fields in the form, and as our system does not depend on authenticated sessions, CSRF tokens do not pose an issue. However, despite considerable active research in breaking CAPTCHAs [11, 60], breaking CAPTCHAs remains out of the scope of this project.

Due to the growing emphasis on responsive web applications, more and more web applications are being built with only client-side JavaScript. Even conventional web applications use JavaScript to dynamically insert content and update the pages. This trend means that these dynamically injected HTML components are not part of the initial HTML that is sent to the client by the server.

Method	Running time	Slowdown
Using our pipeline	629.043	-
Our pipeline with Selenium	919.372	31.58%
Parsing e-mail fields instead of 'grep'ing	707.154	11.05%

Table 7: Running times in seconds for crawling, parsing, and detecting presence of e-mail fields in 1000 random web pages.

Thus, our system will not see dynamically injected forms and hence is unable to detect if E-mail header injection vulnerabilities are present in these forms. The workaround would be to use a JavaScript engine to query for the document.`getElementsByTagName('html')[0].innerHTML` (from inside web browser automation tools such as Selenium), then use that as the source HTML.

A comparison of the running times between the different approaches is shown in Table 7. We chose not to use Selenium as it results in a slowdown of 31.58%.

Because we search for the words e-mail, mail, or email within the HTML form, if the website does not use English names for its forms, our system will not be able to find the presence of an e-mail field. An example is by using the French word for e-mail—courrier électronique—our system is unable to find the presence of the e-mail field.

During the crawl, our system was blacklisted by a few web applications (mostly WordPress ones), and Internet Service Providers (ISPs). To overcome this, we did two things: (1) used an IP range of 60 different IP addresses, and (2) Used a blacklist of our own to prevent our Fuzzer from fuzzing applications that are known to blacklist automated crawlers. This restricted us from gathering data about these applications.

We found that certain WordPress plugins prevent the E-mail header injection attack by sanitizing user input on contact forms. Although not all WordPress web applications are secure, between the presence of the plugins on some websites, and getting tagged as “spambots” by others, we found few vulnerabilities on WordPress web applications.

E-Mail libraries such as the PHP Extension and Application Repository’s (PEAR) mail library provide sanitization for user input. While this is not strictly a limitation of our project, it still means that we are not able to inject sites that used these libraries.

The parser that we use for HTML parsing—Beautiful Soup—does not parse heavily malformed HTML and throws an exception on encountering such HTML. Thus, we have designed the system to exit gracefully on such occasions. A side-effect of this is that our system is unable to test web applications with very bad HTML markup.

Black-box testing is highly beneficial as explained in Section 3.1, however it also has a drawback in that we cannot verify whether the reported vulnerability exists in the source code or is a feature of the website (e.g., the website allows users to send bulk e-mail, adding many cc or bcc headers). We must manually notify the developers to get this feedback.

5.3 Ethics

To make sure that our system did not cause any harm to the web applications that we crawled, we made sure that we did not inject any special characters other than the newline character. We also had an informational website at the IP that we crawled from that described what E-mail header injection was, and contained our

Language	Mail Libraries
PHP	PEAR Mail[26], PHPMailer[44], Swiftmailer[56]
Python	SMTPLib with email.header.Header
Java	Apache Commons E-Mail[13]
Ruby	Ruby Mail >= 2.6[49]

Table 8: Mail libraries that prevent e-mail header injection.

contact details in case the developers of the web applications we crawled wanted to contact us. We maintained a separate blacklist of domain names that the owners did not want us to crawl, and ensured that our system did not crawl their domains. And, the Crawler respected the robots.txt directive of the web application.

5.4 Mitigation Strategy

After demonstrating that E-mail header injection vulnerabilities exist on the web at large, we now describe the most common measures that can be taken to prevent the occurrence of this vulnerability, or at least reduce the impact.

Using a safe, well tested e-mail library is the preferred way of preventing E-mail header injection vulnerabilities (removing the burden of input sanitization from the developer). A list of known secure libraries for each language and framework discussed is shown in Table 8.

Content management systems such as WordPress and Drupal include libraries and plugins to prevent E-mail header injection. Thus, websites built with such CMSes are usually resistant to these attacks. However, it is advised to use the correct e-mail plugin, as not all plugins might be secure.

If neither of the two options are feasible (in-house production, or lack of support infrastructure), developers can choose to perform proper input sanitization. Sanitization should be done with RFC5322 [48] in mind to ensure that all edge cases are covered.

6 RELATED WORK

There are different approaches to finding vulnerabilities in web applications, and most approaches use either black-box testing or white-box testing. Our work is based on the black-box testing approach to finding vulnerabilities on websites, and research has made use of this methodology to find vulnerabilities in web applications [15, 29, 34, 42, 61]. There has been significant discussion on both the benefits of such an approach [14] and its shortcomings [22, 23].

Our work does not intend to act as a vulnerability scanner, but as a means to identify an E-mail header injection vulnerability in a given web application. In this sense, because we are injecting payloads into the web application, our work is related to other injection based attacks, such as SQL Injection [18, 27, 50], Cross-Site Scripting [32, 35], HTTP Header Injection [33], and the related Simple Mail Transfer Protocol (SMTP) Injection [57].

The attack described by Terada [57] is one that attacks the underlying SMTP mail servers by injecting SMTP commands (which are closely related to E-Mail headers and usually have a one-to-one mapping, e.g., To e-mail header has a corresponding To SMTP header) to exploit the SMTP server’s pipelining mechanism. Terada also describes proof-of-concept attacks against certain mailing libraries such as Ruby Mail and JavaMail. This attack, although trying to achieve a similar result, is distinctly different from ours. Terada’s paper also makes this observation and discusses why it is different from E-mail header injection.

In comparison, our work tries to exploit application-level flaws in user input sanitization, which allow this attack. Our work does not intend to exploit the pipelining mechanism, but to exploit the implementation of the mail function in most popular programming languages, which leaves them with no way to distinguish between user supplied headers and headers that are legitimately added by the application.

Although E-mail header injection vulnerabilities have been present for over a decade, there has not been much written about it in the literature, and we find only a few articles on the Internet describing the attack.

The first documented article dates to over a decade ago; a late 2004 article on phpsecure.info [58] accredited to user tobozo describing how this vulnerability existed in the reference implementation of the mail function in PHP, and how it can be exploited. Following this, we found other blog posts [19, 36, 39, 40, 45], each describing how to exploit the vulnerability by using newlines to camouflage headers inside user input. A wiki entry [52] also describes the ways to prevent such an attack. However, none of these articles have performed these attacks against real-life websites.

Another blog post written by user Voxel@Night [3], recounts an actual attack against a WordPress plugin, Contact Form², with a proof of concept. It also showcases the vulnerable code in the plugin that causes the vulnerability. However, this article targets just one plugin and does not aim to find the prevalence of said plugin usage. Neither does it inform the creators of the plugin to fix the discovered vulnerability. This plugin is used actively on 300,000 websites (according to [16]), but is yet to be fixed. E-mail header injection vulnerability was described briefly by Stuttard and Pinto in their book, “*The Web Application Hacker’s Handbook*” [54]. The book, however, does not go into detail on either the attack or the ways to mitigate such an attack. Our work, on the other hand discusses the means to mitigate the attack. We also describe, in detail, the payloads that can be used and the need for varying the payloads (Section 3.2).

To the best of our knowledge, no other research has been conducted to determine the prevalence of this vulnerability on the World Wide Web. We have managed to, on a large scale, crawl and inject web applications with comparatively benign payloads (the bcc header) to identify the existence of this vulnerability without causing any ostensible harm to the website. Our injected payloads *do not contain any special characters other than the newline character* and thus cannot cause any unintended consequences. Also, as we are only injecting one additional bcc header, the underlying mail servers should not be affected by excess load. Our work serves to not only prove the existence of the vulnerability on the World Wide Web but to quantify its prevalence.

7 CONCLUSIONS

We have showcased a novel approach involving black-box testing to identify the presence of E-mail header injection in a web application. Using this approach, we have demonstrated that our system was able to crawl 23,553,796 web pages finding 7,354,425 forms, out of which 1,228,774 forms were fuzzable. We fuzzed 1,012,530 forms and found 74,192 forms that allowed us to send/receive e-mails. Out of these, we were able to inject malicious payloads into 64,510 forms, identifying 994 vulnerable forms (1.54% success rate) across

414 domains. 135 of these domains are found in the Alexa top 1 million websites, and 137 of the vulnerable domains have anti-spoofing mechanisms in place. We also found evidence that 157 IPs that were vulnerable to E-mail header injection are on spamming blacklists. This indicates that the vulnerability is widespread, and needs attention from both web application and library developers.

We hope that our work sheds light on the prevalence of this vulnerability and that it ensures that the implementation of the mail function in popular languages is fixed to differentiate between User-supplied headers, and headers that are legitimately added by the application.

REFERENCES

- [1] Apache Nutch. <http://nutch.apache.org/>
- [2] ContactForm7. <https://wordpress.org/plugins/contact-form-7>
- [3] Vexatious Tendencies. <https://vexatioustendencies.com/wordpress-plugin-vulnerability-dump-part-2/> (2014)
- [4] CVE - Common Vulnerabilities and Exposures (CVE) (2016), <http://cve.mitre.org/>
- [5] ICANN WHOIS Data. <https://whois.icann.org/en> (2016)
- [6] VirusTotal - Free Online Virus, Malware and URL Scanner (2016), <https://www.virustotal.com/>
- [7] Alexa Rankings. <https://data.alexa.com/data?cli=10&url=%URL%> (2017)
- [8] BuiltWith Website Data. <https://builtwith.com/> (2017)
- [9] Wappalyzer. <https://wappalyzer.com/> (2017)
- [10] Acunetix: AcuMonitor: For detecting Email Header Injection, Blind XSS and SSRF - Acunetix. <http://www.acunetix.com/vulnerability-scanner/acumonitor-blind-xss-detection/>
- [11] von Ahn, L., Blum, M., Langford, J.: Telling humans and computers apart automatically. *Commun. ACM* 47(2) (2004), <http://doi.acm.org/10.1145/966389.966390>
- [12] Allman, E., Callas, J., Delany, M., Libbey, M., Fenton, J., Thomas, M.: Domainkeys identified mail (dkim) signatures. *Tech. rep.* (2007)
- [13] Apache Commons Email: (2016), <https://commons.apache.org/proper/commons-email>
- [14] Bau, J., Bursztein, E., Gupta, D., Mitchell, J.: State of the art: Automated black-box web application vulnerability testing. In: *Security and Privacy (SP), 2010 IEEE Symposium on*. pp. 332–345 (May 2010)
- [15] Beizer, B.: *Black-box Testing: Techniques for Functional Testing of Software and Systems*. John Wiley & Sons, Inc., New York, NY, USA (1995)
- [16] BestWebSoft: Contact Form by BestWebSoft WordPress Plugins. <https://wordpress.org/plugins/contact-form-plugin/> (2016)
- [17] Bhide, C.W., Singh, J., Oestreicher, D.: Performance optimizations for computer networks utilizing http (Dec 22 1998), uS Patent 5,852,717
- [18] Boyd, S.W., Keromytis, A.D.: *Sqlrand: Preventing sql injection attacks*. In: *Applied Cryptography and Network Security*. pp. 292–302. Springer (2004)
- [19] Calin, B.: Email Header Injection Web Vulnerability - Acunetix. <https://www.acunetix.com/blog/articles/email-header-injection-web-vulnerability-detection/> (2013)
- [20] Chandramouli, S.P., Zhao, Z., Doupe, A., Ahn, G.J.: E-mail Header Injection Vulnerabilities. *it - Information Technology* 59(2), 67–72 (2017)
- [21] Crocker, D.: Internet Message Format - RFC 2142 (1997), <https://www.ietf.org/rfc/rfc2142>
- [22] Doupe, A., Cavendon, L., Kruegel, C., Vigna, G.: Enemy of the state: A state-aware black-box web vulnerability scanner. In: *Presented as part of the 21st USENIX Security Symposium (USENIX Security 12)*. pp. 523–538. USENIX, Bellevue, WA (2012), <https://www.usenix.org/system/files/conference/usenixsecurity12/sec12-final225.pdf>
- [23] Doupe, A., Cova, M., Vigna, G.: Why johnny can’t pentest: An analysis of black-box web vulnerability scanners. In: *Detection of Intrusions and Malware, and Vulnerability Assessment*, pp. 111–131. Springer (2010)
- [24] Felten, E.W., Balfanz, D., Dean, D., Wallach, D.S.: Web spoofing: An internet con game. *Software World* 28(2) (1997)
- [25] Fielding, R., Gettys, J., Mogul, J., Frystyk, H., Berners-Lee, T., Masinter, L., Leach, P.: RFC 2616 (1999), <https://www.ietf.org/rfc/rfc2616.txt>
- [26] Hagenbuch, C., Heyes, R., Machniak, A.: PEARMail (2016), <https://pear.php.net/package/Mail>
- [27] Halfond, W.G., Viegas, J., Orso, A.: A classification of sql-injection attacks and countermeasures. In: *Proceedings of the IEEE Symposium on Secure Software Engineering* (2006)
- [28] Herzog, A.: Full Disclosure: JavaMail SMTP Header Injection via method setSubject [CSNC-2014-001] (2014), <http://seclists.org/fulldisclosure/2014/May/81>
- [29] Huang, Y.W., Huang, S.K., Lin, T.P., Tsai, C.H.: Web application security assessment by fault injection and behavior monitoring. In: *Proceedings of the 12th International Conference on World Wide Web. WWW ’03, ACM* (2003), <http://doi.acm.org/10.1145/775152.775174>
- [30] Internet Live Stats: www.internetlivestats.com (2016)

²Note that this plugin is used actively on 300,000 websites[16], but is yet to be fixed.

- [31] Jakobsson, M., Myers, S.: Phishing and countermeasures: understanding the increasing problem of electronic identity theft. John Wiley & Sons (2006)
- [32] Jim, T., Swamy, N., Hicks, M.: Defeating script injection attacks with browser-enforced embedded policies. In: Proceedings of the 16th International Conference on World Wide Web. pp. 601–610. WWW '07, ACM, New York, NY, USA (2007), <http://doi.acm.org/10.1145/1242572.1242654>
- [33] Johns, M., Winter, J.: Requestrodeo: Client side protection against session riding. In: Proceedings of the OWASP (2006)
- [34] Kals, S., Kirda, E., Kruegel, C., Jovanovic, N.: Secubat: a web vulnerability scanner. In: Proceedings of the 15th international conference on World Wide Web. pp. 247–256. ACM (2006)
- [35] Klein, A.: [DOM Based Cross Site Scripting or XSS of the Third Kind] Web Security Articles - WebApp Sec (2005), <http://www.webappsec.org/projects/articles/071105.shtml>
- [36] Kohler, D.: damonkohler: Email Injection. <http://www.damonkohler.com/2008/12/email-injection.html> (2008)
- [37] Kucherauw, M., Zwicky, E.: Domain-based message authentication, reporting, and conformance (dmarc) (2015)
- [38] Lin, X., Zavorsky, P., Ruhl, R., Lindskog, D.: Threat modeling for csrf attacks. In: CSE (3). pp. 486–491 (2009)
- [39] Mohamed, A.: PHP Email Injection Example - InfoSec Resources. <http://resources.infosecinstitute.com/email-injection/> (2013)
- [40] Nicol, J.: Securing PHP Contact Forms. <http://jonathannicol.com/blog/2006/12/09/securing-php-contact-forms/> (2006)
- [41] OWASP: https://www.owasp.org/index.php/OWASP_Top_10
- [42] Payet, P., Doupé, A., Kruegel, C., Vigna, G.: EARs in the Wild: Large-Scale Analysis of Execution After Redirect Vulnerabilities. In: Proceedings of the ACM Symposium on Applied Computing (SAC). Coimbra, Portugal (March 2013)
- [43] PHP-Manual: PHP mail - Send mail. <http://php.net/manual/en/function.mail.php> (2016)
- [44] PHPMailer: <https://github.com/PHPMailer/PHPMailer>
- [45] Pope, A.: Prevent Contact Form Spam Email Header Injection | Storm Consultancy Web Design Bath (2008), <https://www.stormconsultancy.co.uk/blog/development/dev-tutorials/secure-your-contact-form-against-spam-email-header-injection/>
- [46] Pope, C., Kaur, K.: Is it human or computer? defending e-commerce with captchas. IT Professional 7(2) (Mar 2005)
- [47] Raghavan, S., Garcia-Molina, H.: Crawling the hidden web. Technical Report 2000-36, Stanford InfoLab (2000), <http://ilpubs.stanford.edu:8090/456/>
- [48] Resnick, P.W.: Internet Message Format - RFC 5322 (2008), <https://tools.ietf.org/html/rfc5322>
- [49] Ruby Mail Gem: <https://rubygems.org/gems/mail>
- [50] Sadeghian, A., Zamani, M., Manaf, A.A.: A taxonomy of sql injection detection and prevention techniques. In: Informatics and Creative Multimedia (ICICM), 2013 International Conference on. pp. 53–56. IEEE (2013)
- [51] Schlitt, W., Wong, M.W.: Sender policy framework (spf) for authorizing use of domains in e-mail, version 1 (2006)
- [52] Email Injection - Secure PHP Wiki. <http://securephpwiki.com/index.php/EmailInjection> (2010)
- [53] Shkapenyuk, V., Suel, T.: Design and implementation of a high-performance distributed web crawler pp. 357–368 (2002)
- [54] Stuttard, D., Pinto, M.: The Web Application Hacker's Handbook: Finding and Exploiting Security Flaws. John Wiley & Sons (2011)
- [55] Su, Z., Wassermann, G.: The essence of command injection attacks in web applications. In: ACM SIGPLAN Notices. vol. 41, pp. 372–382. ACM (2006)
- [56] SwiftMailer: <http://swiftmailer.org/>
- [57] Terada, T.: SMTP Injection via recipient email addresses. MBSD White Paper (December 2015)
- [58] Tobozo: Mail headers injections with PHP. <http://www.phpsecure.info/v2/article/MailHeadersInject.en.php> (2004)
- [59] W3techs: Usage Statistics and Market Share of PHP for Websites, February 2016. <http://w3techs.com/technologies/details/pl-php/all/all> (2016)
- [60] Yan, J., Ahmad, A.S.E.: Breaking visual captchas with naive pattern recognition algorithms. In: Computer Security Applications Conference, ACSAC 2007 (Dec 2007)
- [61] Zanero, S., Carettoni, L., Zanchetta, M.: Automatic detection of web application security flaws. Black Hat Briefings (2005)